

Who Needs a Savings Nudge? Leakage-Aware Prediction of Household Savings Adequacy in India

Rishabh Agrawal, Prisha Kothari, Akshit Kashyap

Goa Institute of Management

Big Data Analytics

Student IDs: {B2025100, B2026092, B2026059}

Email: {rishabh.agrawal25b, prisha.kothari2026b, akshit.kashyap2026b}@gim.ac.in

Abstract—A savings-product or financial-inclusion team cannot review tens of thousands of households by hand to decide who needs support, so the triage has to be automated. We study that triage on the India Human Development Survey-II (41,518 households, 2011–12), predicting whether a household retains at least 20% of its annual income after recorded consumption. The label is an accounting identity over income and total consumption, so expense-to-income ratios and raw rupee categories reconstruct it at 99.75% agreement. Spending composition (each category’s share of total expenditure) looks safer but also leaks: food spending tracks household size, so a formula built from household size, the food share and income ranks households at ROC-AUC 0.895 without fitting any model. We therefore report an income-and-demographics model as the headline and treat spending shares as diagnostic. Under cross-validation grouped by primary sampling unit, the headline XGBoost model reaches ROC-AUC 0.882 and macro-F1 0.782 (0.886 and 0.781 on held-out sampling units), against 0.929 when the shares are added. A single income threshold reaches macro-F1 0.742, and at a 25% contact budget the model reaches 98.5% of the attainable at-risk households against 95.4% for a poorest-first rule, a gain of 1.1 percentage points (95% CI 1.0–1.3). Income carries 58% of the model’s grouped SHAP attribution. Peer benchmarks cannot tell structural from behavioural shortfall: the share of at-risk households who could close their gap by spending like their peers is 22% or 93%, depending on how the benchmark is built.

Index Terms—Machine learning, financial inclusion, household survey data, data leakage, compositional data, explainable AI, grouped cross-validation

I. INTRODUCTION

Savings products in India are sold into a population that is mostly not saving. Outreach is expensive per contact, budgets are fixed, and a team that wants to prioritise households by need has no way to inspect tens of thousands of them individually. The decision is a triage: treat this household as on track or at risk, and spend the next contact accordingly.

We frame that triage as binary classification on the India Human Development Survey-II (IHDS-II), a nationally representative survey of 42,152 Indian households conducted in 2011–12 [1]. IHDS-II is unusual among Indian surveys in recording income and category-level expenditure for the same household; the NSS-lineage consumption surveys collect one without the other. After dropping 634 households with non-

positive income or missing consumption, 41,518 remain, in 2,461 primary sampling units (PSUs).

The problem is constrained in a way that is easy to miss. A household is labelled on track when it retains at least 20% of its annual income, and that label is an arithmetic function of income and total consumption. Expense-to-income ratios, the intuitive representation, reconstruct the answer at 99.75% agreement. Composition shares describe only the shape of a budget and pass every single-feature check, but jointly they are not safe: given household size and income, the food share approximately recovers the size of the budget, and with it the label. The paper is organised around that finding and around a second constraint on reporting: against the rule a team would use without a model, contact the poorest first, most of the apparent gain disappears.

A. Scope and objectives

The project has four objectives: establish by measurement, univariate and multivariate, which representations of household spending can legitimately serve as features for a savings-adequacy target; compare seven model families under PSU-grouped validation fixed in advance and quantify the gain over a single-variable rule rather than over chance; explain the winning model with exact tree SHAP values grouped into feature families and cross-checked by permutation; and translate the results into targeting recommendations that state where the model fails and which spending-based conclusions the data cannot support.

II. RELATED WORK

Predicting household welfare from survey covariates is the core of proxy means testing, the standard method for targeting transfers when income is unobservable. Brown, Ravallion and van de Walle evaluated such tests across nine African countries and found that they miss a large share of the poorest households even when the underlying regression fits well [2], which is the closest antecedent to our finding that a good score does not imply a useful targeting instrument. McBride and Nichols showed that out-of-sample validation and machine learning improve on the regression practice then in use [3], and

where survey data is unavailable poverty has been predicted from mobile phone metadata [4]. Our task differs in what is predicted: the outcome is a savings residual rather than a consumption level, which is what makes spending-based features dangerous.

Ours is a prediction policy problem in the sense of Kleinberg et al. [5]: we do not claim that contacting a household changes its savings rate, only that ranking by predicted risk allocates a fixed budget better than not ranking. The intervention literature supplies the reason such a ranking is worth having, since simple reminders raise savings measurably [6], while Dupas and Robinson show that access constraints rather than preferences explain much of the shortfall among poor households [7].

Two measurement results bear on the data. Deaton and Kozel document the divergence between income and consumption aggregates in Indian surveys [8], which we inherit as 55.9% of households reporting consumption above income. Engel’s law, food’s budget share falling as income rises, appears here at Spearman $\rho = -0.280$.

On method we use gradient-boosted trees [9] with SHAP attributions from the exact tree algorithm [10], [11]. Because the eleven spending shares are a composition, we follow Aitchison’s treatment of the simplex [12], cluster on an isometric log-ratio basis [13] and replace zeros multiplicatively [14]; cluster count is chosen by silhouette [15], Davies–Bouldin and Calinski–Harabasz indices.

III. METHODOLOGY

A. Data and target

The analysis file is built from ICPSR study 36151, dataset DS0002 (Household). IHDS-II records consumption in 52 items under two recall windows, 30-day and annual, so the build annualises as $12 \times \text{monthly} + \text{annual}$ and maps items onto 11 expense categories from the codebook. The 11-category total matches the survey’s own aggregate COTOTAL with correlation 0.9932, median difference zero and 97.7% of households within 1%; a label computed from it agrees with ours for 99.75%. Occupation is the worker type contributing the most workers, ties (18% of households) going to the more regular income stream. A missing outstanding-debt value is kept as missing, with an indicator, rather than read as zero.

The target is

$$\text{Goal_Met} = \mathbf{1} \left[\frac{\text{INCOME} - \text{COTOTAL}}{\text{INCOME}} \geq 0.20 \right], \quad (1)$$

with both quantities in annual rupees. The rate is 31.93% positive, a 2.13:1 imbalance, and 30.47% when survey weights are applied. The 20% figure is a convention, so we publish the sensitivity curve: the at-risk share moves from 68.1% at the benchmark to 61.7% at 10% and 74.7% at 30%, with about 93.5% label agreement either way. Because IHDS records no savings intention, the benchmark is normative and identical for every household; the model measures savings capacity, not goal alignment.

B. The leakage check

Equation (1) is exact over INCOME and COTOTAL, and approximate over the 11 categories. The raw rupee categories, with income, reconstruct the label at 99.75% agreement, and expense-to-income ratios give the same 99.75%. Both are excluded.

Composition shares sum to one, so no single share carries the label: the largest marginal ROC-AUC distance from 0.5 is 0.100 (healthcare share). The multivariate route is different. Food spending rises with household size (elasticity 0.63, $R^2 = 0.29$ in logs), so $\text{COTOTAL} \approx k \cdot \text{size}/\text{food share}$, and the score $1 - (k \cdot \text{size}/\text{food share})/\text{INCOME}$, with k the median food spend per person (Rs. 8,848 a year), ranks households at ROC-AUC 0.895 with no fitted model. It excludes only the 24 households (0.06%) with a zero food share. In an ablation under grouped cross-validation, dropping the food share alone costs 0.002 AUC (the other shares substitute for it), dropping the household-size family costs 0.010, and dropping all shares and participation indicators keeps 94.3% of the full model’s AUC (0.878 against 0.931). We fixed a decision rule before running these checks: if the formula scored above 0.85, or a model without shares kept most of the full model’s AUC, the headline model would exclude spending shares. Both conditions hold.

Six survey-reported indicators of savings-instrument holding sit outside the consumption arithmetic and are reserved as an external check: fixed-deposit holding is 15.2% among on-track households against 8.2% among the rest.

C. Features and validation

Two feature sets are carried. The *deployable* set (13 features) is log income, household size, age dependents (members aged 0–14 or 60+), dependency ratio, head age, highest adult education, debt-to-income, a debt indicator, a missing-debt indicator and four one-hot categoricals (occupation, area type, caste, religion) with an explicit Unknown level. The *full* set (29) adds the 11 shares and five participation indicators for majority-zero categories.

All design decisions were taken on a training split drawn first: one fold of StratifiedGroupKFold grouped by PSU, leaving 8,299 test households in 472 PSUs that no decision saw. Everything else uses five-fold stratified cross-validation grouped by PSU on the remaining 33,219 households, because neighbours share prices, labour markets and interviewers. Debt is winsorised at the 99th percentile inside the pipeline, refit per fold. Linear models use standard scaling and drop a reference share (food) so the remaining coefficients are identified; tree models are unscaled. Robust scaling was dropped because zero-inflated shares have an interquartile range near zero, which inflated entertainment shares to values above 150.

The feature checks behave as before: expense-to-income ratios transfer worst across income halves on a neutral target (Has_Bank_Savings) even after winsorising (mean ROC-AUC 0.603 against 0.628 for raw rupees and 0.618 for shares); participation indicators lift ROC-AUC from 0.917 to 0.919;




Fig. 1. Grouped-CV ROC-AUC on the training split as features and capacity are added. The red line is the score from household size, food share and income, with no fitted model: it beats every model that does not see spending shares, and the full model’s extra 0.047 AUC over the headline runs largely through that route.

TABLE I

FIVE-FOLD PSU-GROUPED CROSS-VALIDATION ON THE TRAINING SPLIT (33,219 HOUSEHOLDS), ALL ROWS INCLUDING THE BASELINES. XGBOOST (TUNED) USES THE CONFIGURATION IN RESULTS/MODEL_FINAL.JSON. “FULL” ADDS THE SPENDING SHARES AND IS DIAGNOSTIC ONLY.

Model	Deployable (headline)				Full AUC
	F1	AUC	PR-AUC	ECE	
XGBoost (tuned)	0.782	0.882	0.942	0.011	0.929
HistGradientBoosting	0.778	0.880	0.941	0.012	0.928
XGBoost (untuned)	0.774	0.875	0.938	0.023	0.929
Logistic Regression	0.773	0.876	0.938	0.103	0.920
Linear SVM	0.774	0.876	0.938	—	0.920
Random Forest	0.777	0.870	0.934	0.082	0.915
Decision Tree	0.751	0.850	0.914	0.104	0.882
Single income threshold	0.742	0.739	0.808	0.003	0.739
Majority class	0.405	0.500	0.681	0.319	0.500

and a centred log-ratio transform of the shares loses (0.911 against 0.920). A missing debt value is informative (goal attainment 0.367 against 0.315), which the indicator now captures.

Macro-F1 remains the operational ranking metric, but every model is also scored on threshold-free metrics: ROC-AUC, PR-AUC for the at-risk class (positive = 1 − Goal_Met, prevalence 0.681), Brier score, log-loss and expected calibration error (ECE).

IV. RESULTS

A. What each ingredient is worth

The majority classifier reaches 0.681 accuracy and 0.405 macro-F1. A depth-1 tree on income, re-learned inside every training fold (threshold Rs.110,377–126,545, mean Rs. 122,249 a year), reaches 0.778 accuracy and 0.742 macro-F1. Logistic regression on income alone reaches ROC-AUC 0.835, and 0.876 with demographics and debt. Adding the food share to income lifts it to 0.875, and to 0.894 once household size is also present, which is the reconstruction route at work (Fig. 1).

B. Model comparison

On the deployable set the tuned XGBoost model (max_depth=4, learning_rate=0.08,




Fig. 2. SHAP for the headline model on held-out PSUs: the ten largest mean absolute values, and attribution by feature family, computed by summing SHAP values within each family before averaging.

n_estimators=200, subsample=0.7, min_child_weight=1) leads at macro-F1 0.782 and ROC-AUC 0.882 (Table I); histogram gradient boosting is within one fold standard deviation. Tuning was worth 0.007 macro-F1, and depth was the parameter that mattered. PR-AUC is reported for the at-risk class against its 0.681 base rate. The linear models rank close on ROC-AUC but are poorly calibrated (ECE 0.10) because of balanced class weights.

On the held-out PSUs, touched once, the headline model scores ROC-AUC 0.886, macro-F1 0.781, at-risk PR-AUC 0.944, Brier 0.127 and ECE 0.009; the full model scores 0.932 and 0.836. A household-level random split gives almost the same full-model score (0.931), so the gap between the two models comes from the feature set, not from the validation scheme.

For deployment the at-risk threshold is tuned on out-of-fold training scores, not accepted at 0.5. The F1-optimal threshold, applied once to the held-out set, gives at-risk precision 0.826 and recall 0.913. Out of fold, raising at-risk recall from 80% to 95% costs 10 points of precision (0.893 to 0.789).

C. What the model uses

We attribute with shap.TreeExplainer, verified to reproduce the raw margin to 6.5×10^{-6} . Because SHAP values are additive, a family’s effect is the sum of its columns’ values; we average its absolute value, which avoids double counting columns that push in opposite directions. Income carries 58.0% of grouped attribution in the headline model, the household-size family 14.5%, the social and geographic categoricals 10.9%, debt 8.7% and head age with education 8.0% (Fig. 2). Grouped permutation importance agrees on the order and is more extreme: permuting income costs 0.385 ROC-AUC, 82% of the total drop. Interactions are 31.6% of total attribution over the whole held-out set, led by size with income.

In the full model the shares family carries 19.1% of grouped SHAP (16.0% of permutation importance), and the strongest interaction is food share with income. Read naively, this says that a food-dominated budget signals adequacy once income




Fig. 3. Goal attainment by spending cluster within income decile (all ten). Cluster membership adds explanatory power beyond income (likelihood-ratio test), but the clusters are defined by which categories are absent.

is controlled. The mechanism is the reconstruction route: at a given size and income, a higher food share implies a smaller total budget. The food-share SHAP values rise as the implied budget-to-income ratio falls (Spearman 0.27).

The shares are also not identified in a linear model. Refitting on five disjoint subsamples, three of ten share coefficients flip sign with all eleven shares included and one with food dropped as reference. With standard scaling, dropping the reference changes no coefficient’s spread by more than a factor of 1.4 in either direction. Robust scaling is avoided because a zero-inflated share such as rent has an interquartile range of zero.

D. Spending clusters

We clustered on one representation only, five isometric log-ratio coordinates of the six least zero-inflated shares, and chose k within it. Silhouette (0.412), Davies–Bouldin (0.98) and Calinski–Harabasz all select $k = 3$, and the partition is stable under bootstrap resampling (median adjusted Rand index 0.998). It is not, however, a set of behavioural personas. The clusters are households with no transport spend (11.7%), with no healthcare spend (18.6%), and the rest (69.7%); their adjusted Rand index with the raw zero pattern is 0.858. Varying the zero-replacement constant from 0.1 to 1 times the smallest positive share keeps the partition (index 0.93–1.00 against the baseline) but moves the silhouette that justifies it from 0.39 to 0.48.

The clusters are close to independent of income (normalised mutual information with income decile 0.012; Kruskal–Wallis $\epsilon^2 = 0.034$). Unconditionally their goal attainment differs by 9.2 points (Cramér’s $V = 0.076$). Conditioning on income decile, cluster membership is significant (likelihood ratio 774 on 2 df; pseudo- R^2 0.265 to 0.280), with an income-adjusted odds ratio of 0.34 for the spend-on-everything cluster against the no-transport cluster, and the effect varies by decile (interaction $p = 10^{-19}$; Fig. 3). The mechanism is the reconstruction route again: a household with an empty category spends less in total at a given income.

E. What the model is worth

At a 25% contact budget no strategy can reach more than 36.7% of at-risk households, because they are 68.1% of the population. The headline model reaches 36.2% (98.5% of the ceiling) and a poorest-first income rule 35.0% (95.4%; Fig. 4(a)). The PSU-bootstrap interval on the difference is 1.0 to 1.3 percentage points; at 10%, 50% and 65% budgets it is 0.3, 2.7 and 2.7 points. Against the income rule, not against random contact, the model’s precision at 25% is 3.1 points higher (98.5% against 95.4%). Out-of-fold calibration is good (ECE 0.006).

Accuracy is lowest where the decision is contested: 0.98 in the poorest decile, 0.83 in the richest, and 0.66 in decile 7 (Fig. 4(b)), where at-risk rates are near one half.

To look for recoverable spending we compare each at-risk household with the mean on-track household in its income decile, with PSU-bootstrap intervals. In rupees at-risk households out-spend their peers in every category, which is mechanical: at a given income, being at risk means a larger budget. Composition is more informative (Fig. 4(c)). At-risk budgets give more to healthcare (+4.4 points, CI 3.8–4.9), education (+2.8), miscellaneous (+2.1, CI 1.7–2.5) and transport (+1.8), and less to food (−9.0). Miscellaneous spending is third, and the two largest categories are not legitimate nudge targets. Whether at-risk households could close their gap by spending like their peers depends on the benchmark: 92.8% could by matching peers’ rupee spending where it is lower, 22.3% by matching their budget shares. IHDS therefore cannot say what fraction of the shortfall is structural.

Ranking by a quantile model’s predicted rupee shortfall, instead of a probability, reaches 45.5% of the total rupee gap at a 25% budget against 37.0% for the classifier and 28.7% for the income rule.

Five recommendations follow, each computed in the notebook that produces it: prioritise by income and use the onboarding model to re-rank within income bands; judge targeting against the ceiling and the income rule rather than random contact; do not size a budgeting product from peer benchmarks; treat miscellaneous spend as one hypothesis among several rather than the lever; and report relative priority only, never a prevalence figure.

V. DISCUSSION

A model trained on spending shares scores 0.93 ROC-AUC, but a three-column formula with no fitted parameters beyond a median scores 0.895, and removing the shares costs only 0.047. The model we recommend uses what a team would know at onboarding: income, household composition, debt and social group. It reaches 0.886 on held-out sampling units and improves on a poorest-first rule by about one percentage point of attainable capture at a 25% budget, with an interval that excludes zero. That gain is small, consistent with Brown et al.’s experience of proxy means tests [2].

Three general lessons follow. First, a feature can pass every univariate leakage check and still reconstruct the label jointly




Fig. 4. Three views of deployment value, all out-of-fold under PSU-grouped CV. (a) At-risk households reached against contact budget, with the ceiling $\min(1, b/0.681)$. (b) Accuracy of the headline model by income decile. (c) Signed mean difference in budget share between at-risk households and on-track households in the same income decile, with 95% PSU-bootstrap intervals.

with others; composition shares do so whenever one category is close to proportional to a measured covariate. Spending shares belong in a deployed model only if they are observed before the prediction window, which is how the scoring service in the repository enforces them. Second, capture should be reported against its ceiling and against the obvious rule; with a majority at-risk class, lift over random is uninformative. Third, peer-benchmark findings about “recoverable” spending are only as good as the benchmark, and here two reasonable constructions disagree by a factor of four.

Income under-reporting remains the largest limitation. Consumption exceeds income for 55.9% of households (57.7% weighted), 22.0% report spending more than twice their income, and all of the latter are at risk, 32.3% of that group. This is a documented property of Indian surveys [8]; relative rankings hold and absolute prevalence claims do not. The data is from 2011–12; the unit is a household; model fitting is unweighted while population figures use the survey weights; debt is a stock with no repayment flow; and the 20% benchmark is a convention. IHDS records no intervention, so nothing here estimates who would respond to a nudge; an uplift model needs campaign data.

VI. REPRODUCIBILITY

Code, walkthrough documents and result artifacts are at <https://github.com/saint2706/savings-goal-classifier>. ICPSR’s terms prohibit redistributing the IHDS-II microdata, so `dataset/README.md` documents the free registration and download. From a clean clone: `uv sync`, then `uv run sgc build --tsv <path>/36151-0002-Data.tsv` (the path is required), then the eight notebooks in numeric order, about 25 minutes on four cores. The build is validated against a Pandera schema; the package is covered by tests on a synthetic survey file, run in CI with `ruff` and `mypy`. Every random seed is 42, every split is grouped by PSU, and the figures here are regenerated by `project/figures/make_figures.py` from the result files the notebooks write.

A. AI Use Declaration

Claude Code (Anthropic), an AI coding assistant, was used to write the analysis package, notebooks, tests and figure script, and to draft this report’s text from the resulting artifacts. We verified the numbers against the committed result files and revised the text. The analysis design and interpretation are our own, as recorded in the walkthroughs. We make no claim about the proportion of AI-generated text.

VII. CONCLUSION

We built a savings-adequacy classifier for 41,518 Indian households and found that spending composition, the representation usually chosen to avoid leakage, leaks through household size. The model we recommend uses income, demographics and debt, reaches ROC-AUC 0.886 on held-out sampling units, and adds about a percentage point of attainable at-risk capture over contacting the poorest first. It belongs inside an income-ranked campaign, not in front of one.

The broader result is methodological. A normative savings threshold turns most spending-based features into restatements of the label, singly or in combination, and conclusions drawn from them need a multivariate leakage check, grouped validation and uncertainty intervals before they can be trusted. Here those checks rule out a reversed Engel effect, a single miscellaneous-spending lever, and any firm split between structural and behavioural shortfall.

ACKNOWLEDGMENT

We thank our course instructor for feedback on the framing of this project, and NCAER and ICPSR for collecting and distributing the IHDS-II microdata.

REFERENCES

- [1] S. Desai, R. Vanneman, and NCAER, “India Human Development Survey-II (IHDS-II), 2011–12,” ICPSR study 36151, 2018. doi: 10.3886/ICPSR36151.v6
- [2] C. Brown, M. Ravallion, and D. van de Walle, “A poor means test? Econometric targeting in Africa,” *J. Develop. Econ.*, vol. 134, pp. 109–124, 2018.

- [3] L. McBride and A. Nichols, "Retooling poverty targeting using out-of-sample validation and machine learning," *World Bank Econ. Rev.*, vol. 32, no. 3, pp. 531–550, 2018.
- [4] J. Blumenstock, G. Cadamuro, and R. On, "Predicting poverty and wealth from mobile phone metadata," *Science*, vol. 350, no. 6264, pp. 1073–1076, 2015.
- [5] J. Kleinberg, J. Ludwig, S. Mullainathan, and Z. Obermeyer, "Prediction policy problems," *Amer. Econ. Rev.*, vol. 105, no. 5, pp. 491–495, 2015.
- [6] D. Karlan, M. McConnell, S. Mullainathan, and J. Zinman, "Getting to the top of mind: how reminders increase saving," *Manage. Sci.*, vol. 62, no. 12, pp. 3393–3411, 2016.
- [7] P. Dupas and J. Robinson, "Why don't the poor save more? Evidence from health savings experiments," *Amer. Econ. Rev.*, vol. 103, no. 4, pp. 1138–1171, 2013.
- [8] A. Deaton and V. Kozel, "Data and dogma: the great Indian poverty debate," *World Bank Res. Observer*, vol. 20, no. 2, pp. 177–199, 2005.
- [9] T. Chen and C. Guestrin, "XGBoost: a scalable tree boosting system," in *Proc. ACM SIGKDD*, 2016, pp. 785–794.
- [10] S. M. Lundberg and S.-I. Lee, "A unified approach to interpreting model predictions," in *Proc. NeurIPS*, 2017, pp. 4765–4774.
- [11] S. M. Lundberg et al., "From local explanations to global understanding with explainable AI for trees," *Nature Mach. Intell.*, vol. 2, no. 1, pp. 56–67, 2020.
- [12] J. Aitchison, "The statistical analysis of compositional data," *J. Roy. Statist. Soc. B*, vol. 44, no. 2, pp. 139–177, 1982.
- [13] J. J. Egozcue, V. Pawlowsky-Glahn, G. Mateu-Figueras, and C. Barceló-Vidal, "Isometric logratio transformations for compositional data analysis," *Math. Geol.*, vol. 35, no. 3, pp. 279–300, 2003.
- [14] J. A. Martín-Fernández, C. Barceló-Vidal, and V. Pawlowsky-Glahn, "Dealing with zeros and missing values in compositional data sets using nonparametric imputation," *Math. Geol.*, vol. 35, no. 3, pp. 253–278, 2003.
- [15] P. J. Rousseeuw, "Silhouettes: a graphical aid to the interpretation and validation of cluster analysis," *J. Comput. Appl. Math.*, vol. 20, pp. 53–65, 1987.